

Hadoop Developer Hadoop Developer Hadoop Developer - Liberty Mutual Insurance Dover, NH 8+ years of experience in complete Software Development Life Cycle including analysis, design, administration, development, deployment and maintenance of critical software and Big Data applications. Experience with Apache Hadoop components like HDFS, MapReduce, HiveQL, HBase, Pig, Hive, Sqoop, Oozie, Cassandra, Flume, and Spark. Good knowledge of Hadoop Architecture and various components such as HDFS, Job Tracker, Task Tracker, Name Node, Data Node, YARN and MapReduce concepts. Experience in working with different kind of MapReduce programs using Hadoop for working with Big Data analysis. Experience with Oozie Workflow Engine in running workflow jobs with actions that run Hadoop MapReduce and Pig jobs. Well versed in installation, configuration, supporting and managing of Big Data and underlying infrastructure of Hadoop Cluster. Skilled in Apache Spark jobs using Scala in test environment for faster data processing and used Spark SQL for querying. Knowledge on implementing Big Data in Amazon Elastic MapReduce (Amazon EMR) for processing, managing Hadoop framework dynamically scalable Amazon EC2 instances. Experience in creating Business Intelligence solutions and ETL workflows using Tableau. Experience in data processing like collecting, aggregating, moving from various sources using Apache Flume and Kafka. Hands on experience in developing the applications with Java, J2EE, Servlets, JSP, EJB, SOAP, JNDI, JMS, JDBC2, Hibernate, Struts, Spring, XML, HTML, XSD, XSLT. Experience in writing SQL, PL/SQL queries for accessing and managing databases such as Oracle, SQL Server and NoSQL databases like HBase, Cassandra, MongoDB. Experience in importing and exporting data using Sqoop from Relational Database Systems to HDFS and vice-versa. Experience in configuring the Zookeeper to coordinate servers in clusters and to maintain data consistency. Expertise in loading the data from the different data sources like into HDFS using Sqoop and load into partitioned Hive tables. Good communication, interpersonal and analytical skills and a highly motivated team player with the ability to work independently. Work Experience Hadoop Developer Liberty Mutual Insurance - Dover, NH March 2018 to Present Responsibilities: With more than 211,300 people in over 40 countries, Capgemini is one of the world's foremost providers of consulting, technology and outsourcing

services. The Group reported 2018 global revenues of EUR 13.2 billion. Together with its clients, Capgemini creates and delivers business, technology and digital solutions that fit their needs, enabling them to achieve innovation and competitiveness. A deeply multicultural organization, Capgemini has developed its own way of working, the Collaborative Business Experience(TM), and draws on Rightshore , its worldwide delivery model. Rightshore is a trademark belonging to Capgemini

Worked with business partners in discussing the requirements for new projects and enhancements to the existing applications. Worked with application teams to install operating system, Hadoop updates, patches, version upgrades as required. Developed MapReduce programs to parse the raw data, populate tables and store the refined data in partitioned tables. Used Spark API over Cloudera Hadoop YARN to perform analytics on data in Hive. Involved in creating Hive Tables, loading with data and writing Hive queries which will invoke and run MapReduce jobs in the backend. Developed Pig Latin Scripts to process data in a batch to perform trend analysis. Wrote ETL jobs to read from web APIs using REST and HTTP calls and loaded into Cassandra using Java and Talend. Analyzed Cassandra database and compared it with other open-source NoSQL databases to find which one of them better suites the current requirement. Responsible for Cluster maintenance, Monitoring; commissioning/decommissioning Data Nodes; troubleshooting, manage and reviewing Data Backups and Log Files. Ingested data from RDBMS and performed data transformations, and then export the transformed data to Cassandra as per the business requirement. Developed automated processes for flattening the upstream data from Cassandra which in JSON format. Used Hive UDFs to flatten the JSON Data. Configured Spark Streaming to receive real time data from the Kafka and store the stream data to the database. Responsible for developing data pipeline by implementing Kafka producers and consumers. Connected to AWS EC2 using SSH and ran Spark-submit jobs. Responsible for developing data pipeline with AWS to extract the data from weblogs and store in the database. Analyzed user request patterns and implemented various performance optimization measures including implementing partitions and buckets in HiveQL. Studied data by performing HiveQL & running Pig Latin scripts to study customer behavior. Environment: Hadoop, Hive, MapReduce, Pig

Latin, REST, Java, Cassandra, JSON, Spark, AWS, EC2, HiveQL, Oozie Hadoop Developer  
NorthShore Medical Group - Lincolnwood, IL November 2015 to March 2018 Responsibilities:  
Hadoop Developer able to work in one or more than one projects in the hadoop data lake with  
added expertise in NLP using Python, including technical deliverables as per business needs.  
Hadoop Developer with NLP and Data / text mining skills able to work in one or more than one  
projects in the hadoop data lake, including technical deliverables as per business needs. Essential  
Duties and Responsibilities: Following is a summary of the essential functions for this job. Other  
duties may be performed, both major and minor, which are not mentioned below. Specific activities  
may change from time to time. Worked on a product team using Agile/SCRUM methodology to  
develop, deploy and support solutions that leverage the Client big data platform Documented the  
systems processes and procedures for future references. Supported technical team in  
management and review of Hadoop log files and data backups. Involved in writing MapReduce  
program and Hive queries to load and process data in Hadoop File System. Involved in creating  
Hive tables, loading with data and writing hive queries which will run internally in MapReduce way.  
Developing and running MapReduce jobs on YARN and Hadoop clusters to produce daily and  
monthly reports as per user's need. Maintain and schedule periodic jobs which range from  
updates on MapReduce jobs to creating ad-hoc jobs for the business users. Scheduling and  
managing jobs on a Hadoop cluster using Oozie work flow. Handled importing data from different  
data sources into HDFS using Sqoop and performing transformations using Hive, MapReduce and  
then loading data into HDFS. Exporting of result set from HIVE to MySQL using Sqoop export tool  
for further processing. Collecting and aggregating large amounts of log data using Flume and  
staging data in HDFS for further analysis. Developed several REST web services supporting both  
XML and JSON to perform task such as demand response management. Created Maven builds to  
build and deploy Spring Boot microservices to internal enterprise Docker registry. Designed target  
tables as per the requirement from the reporting team and designed Extraction, Transformation and  
Loading using Talend. Implemented File Transfer Protocol operations using Talend Studio to  
transfer files in between network folders. Optimized MapReduce Jobs to use HDFS efficiently by

using various compression mechanisms. Responsible for continuous monitoring and managing Elastic MapReduce (EMR) cluster through AWS console. Executed speedy reviews and first mover advantages by using workflows like Oozie in order to automate the data. Involved in loading data from UNIX file system to HDFS. Environment: Hadoop, MapReduce, HDFS, Hive, Oracle, Java, AWS, Servlets, HTML, XML, SQL, J2EE, JUnit, Tomcat, UNIX, Maven, REST Hadoop Developer Barnes & Noble - Brentwood, TN March 2014 to September 2015 Responsibilities: Key skills required for the job are: n Hadoop-L3, (Mandatory) .As a Senior Developer, you are responsible for development, support, maintenance and implementation of a complex project module. You should have good experience in application of standard software development principles. You should be able to work as an independent team member, capable of applying judgment to plan and execute your tasks. You should have in-depth knowledge of at least one development technology/ programming language. You should be able to respond to technical queries / requests from team members and customers. You should be able to coach, guide and mentor junior members in the team. Minimum work experience: 3 - 5 YEARS Involved in developing Hadoop Map Reduce jobs using Java Runtime Environment for the batch processing to search and match the scores. Executed speedy reviews and first mover advantages by using workflows like Oozie in order to automate the data. Loading process into the Hadoop distributed File System (HDFS) and Pig language in order to preprocess the data. Integrated Oozie with the rest of the Hadoop stack supporting several types of Hadoop jobs out of the box such as Map-Reduce, Pig, Hive, Sqoop, Flume. Written Hive jobs to parse the logs and structure them in tabular format to facilitate effective querying on the log data. Worked with Sqoop to import export data from HDFS to Relational Database system. Performed performance tuning and troubleshooting of Map Reduce jobs by analyzing and reviewing Hadoop log files. Developed Pig scripts for analyzing large data sets in the HDFS. Collecting and aggregating large amounts of log data using Apache Flume and staging data in HDFS for further analysis. Designed and built the Reporting Application, which uses the Spark SQL to fetch and generate reports on HBase table data. Extracted data from the server into HDFS and Bulk Loaded the cleaned data into HBase.

Extracted data from the flat files and other RDBMS databases into staging area and populated onto Data warehouse. Worked on JVM performance tuning to improve Map-Reduce jobs performance.

Developed programs to manipulate data & perform CRUD operations on request to database. Developed several REST web services supporting both XML and JSON to perform task such as demand response management. Used message driven beans for asynchronous processing alerts to the customer. Worked on developing Use Cases, Class Diagrams, Sequence diagrams, and Data Models. Performed performance tuning and troubleshooting of Map Reduce jobs by analyzing and reviewing Hadoop log files. Environment: Hadoop, Hive, HBase, Linux, MapReduce, HDFS, Hive, Java (JDK), Cloudera, MapReduce, DataStax, IBM DataStage, UNIX Shell Scripting.

Java Developer Charter Communications - St. Louis, MO April 2011 to February 2014

Responsibilities: Around the world, an important revenue driver for the firm. In FICC Technology we use both open source industry standard and internal proprietary technologies to build cutting edge platforms for pricing, execution, and control over each of these millions of transactions. The Fixed Income, Currencies and Commodities (FICC) Technology group develops technology to enable trading across a diverse set of businesses under FICC within Goldman Sachs' Securities Division. As such, it plays a key role in the firm's ability to help institutional and corporate clients trade a variety of financial instruments on electronic venues and over-the-counter. Interacted with business analyst to understand the requirements to ensure correct modules been built to meet business requirements. Participated in the daily SCRUM meetings to produce quality enhancements within time. Developed UML using Case diagrams, Class diagrams, and Sequence diagrams using Rational Software Architect. Spring MVC model integration for front-end request action controller. Developed web screens in JSP, JSTL, CSS and client-side validation using jQuery. Developed Web services to allow communication between application through SOAP over HTTP using Apache CXF. Configured JMS on Web Sphere Server for asynchronous messaging through implementation of Message Driven Beans (MDB). Used Spring ORM module for integration with Hibernate for persistence layer. Implemented the application using the concrete principles laid down by several design patterns such as Session Fa ade, Business Delegate, Singleton, Data

Access Object, and Service Locator. Developed the application in J2EE Application Server environment with IBM WebSphere as deployment server with RAD as development IDE. Used JIRA for defect tracking and project management. Developed and designed XML Schemas to transport and store data. XML was used to simplify data and allow for Platform Changes, as well as making data more available across the applications distributed platforms. Extensively used XSLT to transform XML documents to HTML. Wrote custom jQuery plugins and developed JavaScript functions to build a bleeding-edge, AJAX-driven user interface. Developed unit and functional test cases using J-Unit. Maven and Jenkins used for the automatic build process. Used Log4J utility to log error, info and debug messages. Used Rational Clear Case for version controlling. Worked efficiently in a very tight schedule to meet the deadlines Environment: Java, J2EE, Spring, MVC, Hibernate, HTML, CSS, AJAX, jQuery, JavaScript, JIRA, XML, J-Unit, Maven, Jenkins, Log4J Education Bachelor's Skills AJAX (2 years), Apache (4 years), APACHE HADOOP HDFS (3 years), APACHE HADOOP OOZIE (5 years), APACHE HADOOP SQOOP (3 years), databases (2 years), Flume (3 years), Hadoop distributed File System (3 years), HDFS (3 years), Hive. (5 years), Java (8 years), JavaScript (2 years), jQuery. (2 years), JSON (5 years), Oozie (5 years), REST (5 years), SQL (3 years), Sqoop (3 years), UNIX (3 years), XML (6 years) Additional Information Technical Skills: Big Data HDFS, Map Reduce, Hive, Pig, Sqoop, Flume, Oozie, Zookeeper, Kafka, Impala, Apache Spark, Hue, Ambari, Apache ignite. Languages Java, SQL, PYTHON, PL/SQL, HQL, Scala Web Technologies HTML5, CSS3, JavaScript, jQuery, AJAX, Servlets, JSP, JSON, XML, XHTML, JSF, Angular, SOAP, REST, WSDL, JAXB, and JAXP Operating Systems Windows (XP,7,8), UNIX, LINUX, Ubuntu, CentOS Databases Oracle, MySQL, DB2, HBase, Cassandra

Name: Ashley Dunn

Email: edwardselizabeth@example.org

Phone: 605.296.4354x57294