

Site Reliability Engineer Site Reliability Engineer Site Reliability Engineer Bridgewater, NJ
Authorized to work in the US for any employer Work Experience Site Reliability Engineer
Northwestern Mutual - New York, NY October 2018 to June 2019 Deployed Cassandra cluster in
AWS using Kubernetes Production Support and Deployments on Cassandra Clusters. Involved
in troubleshooting Cassandra performance related issues by analyzing various metrics available
from OpsCenter & Custom Shell Scripts (Read Request Latency, Write Request Latency, Cfstats,
TPstats, netstats, Dstats) Actively Work with Development team and Operation team during crisis
time to minimize customer impact and verifying the data consistency across all datacenter using
CQL (Cassandra Query Language). Data Engineer, ETL/Big Data/PL/Sql/Cassandra/ Python
JPMorgan CHASE & Co - Jersey City, NJ March 2018 to August 2018 The Consumer Bank CRM
will provide an enhanced platform leveraging best of breed Salesforce capabilities aligned to a new
branch relationship building process and branch lead optimization. The strategic platform will
provide features and functionality for retail bankers to manage their book of work, service customers
and deepen customer relationship. Involved in Cassandra Resiliency Architecture, Cassandra
data modelling of Customer Keyspace, Account Keyspace and Lead Keyspace. Production
Support and Deployments on Cassandra Clusters. Involved in supporting Multiple production
Cassandra Clusters with 6 Nodes on each 2 Clusters. Verizon Wireless, Piscataway, NJ
Sys-Eng/MTS-III- Sys-Eng Data Engineer MTS-IV May 2015 to December 2017 Involved in
Cassandra Architecture, Cassandra data modelling & experience in installation, configuration and
monitoring Multiple DataStax Enterprise Cassandra cluster (B2B & B2C). Worked with Datastax
Support team in implementing patches/Upgrades. Involved in supporting Multiple production
Cassandra Clusters with 300 Nodes, 400TB of data across two Clusters: B2C 180 Nodes & B2B
120 Nodes. Involved in end to end data modeling from RDBMS (IBM-DB2) to Cassandra and
actively worked on designing data migration process from DB2 to Cassandra for Read and Write
activities. Actively involved in working with Development and Operations team's in migrate API's
from IBM-DB2 to Cassandra which resulted in significant cost saving to organization. Involved in
evaluating and setting up Cassandra Databases on AWS and bare metal systems for

non-production environment. Lead automation effort in maintaining 300 Nodes Production Cassandra clusters for (Start/Stop/ Upgrade/Repair/Compaction/Monitoring) and worked on User defined compaction to release tombstones in Cassandra. Lead end to end production monitoring, notification setup for Cassandra Clusters using OpsCenter, Node tool utilities, Shell Scripting and Python. Actively worked on troubleshooting hardware ware failures; analyzing disk failures, Memory Leak. etc Involved in troubleshooting Cassandra performance related issues by analyzing various metrics available from OpsCenter & Custom Shell Scripts (Read Request Latency, Write Request Latency, Cfstats, TPstats, netstats, Dstats) Actively Involved in expanding clusters like adding New Datacenter's to existing clusters, Node addition and removing Dead Nodes. Involved in 180 Node cluster upgrades from 4.8.9 to 4.8.14 to resolve significant SStable corruption issues and automating corruption clean up using (Offline Scrub and Online Scrub). Actively Work with Development team and Operation team during crisis time to minimize customer impact and verifying the data consistency across all datacenter using CQL (Cassandra Query Language). Create new Colum Family's as per application migration requirements Support Iconic Launch (IPhone/Samsung Launch) schedules. 24*7 Production support & On- call Support Experience in benchmarking Cassandra Cluster using Cassandra stress tool. Oracle Pl/Sql Developer Verizon Wireless - Township of Warren, NJ August 2011 to May 2015 Executed and review Exacheck results and implement necessary fixes. Work with Oracle Platinum Support team to seamlessly install Quarterly patches. Installed and configured Shareplex replication between multiple datacenters, setup and implement Oracle Change Data Capture (CDC), Oracle AQ/DQ to provide high availability. Involved in overall Application performance optimization by reviewing SQL Execution plans, PL-SQL procedures/packages. Write scripts for creation of users, creation and assignment of roles, analyze space requirements and maintain table spaces, creation and maintenance of tables, partition tables and indexes Stats Collection. Scheduled Jobs for auto drop partitions. Involved in maintaining IBM Change Data Capture like creating Subscriptions, add/drop tables from subscriptions Created Users, Roles, Privileges, Profiles, managing security, analyzing and monitoring the system performance and tune accordingly. Involved in database Export/Import

and performing backups using RMAN. Developed Python code to extract all image URLs from the vendor files, download images and stored in Cassandra for Web-facing Media store Applications. Developed Automated process to retrieve and upload all images to Akamai web server as backup using Python. Support Organizational enterprise release. Production Support (Operations) , Data Engineer - ETL Verizon Wireless - Township of Warren, NJ January 2011 to August 2011 Analyzing and proactively managing problems and making decisions. Validations script to enhance MCS / VISION Grid data compares and Recon and optimized various Oracle Database Objects such as, user defined functions and PL-SQL procedures, Data Import/Export, triggers and packages using Explain plan and hints and various sql loader control files were written. Incident Reporting into JIRA ticketing system and providing necessary analysis for bug fixes and adhering to published procedures and maintain support knowledge base and documentation and contribute to team's increased ability to serve business. Worked with Analyst/business users, Dev team, Infrastructure team (Hardware and capacity planning) within MCS for resolving production bugs, optimizing and tuning reporting Queries. Participated in UAT/training sessions to educate business users about new release functionalities. Used Oracle CDC Replication feature to set up replication from database to grid. Experienced in documentation, provided documentation for database design, analysis, modifications, creating schemas etc. Planning, sizing, scheduling including implementation activities for the Operations/Production Support team and closely with Development, Infrastructure, DBA, Middleware and Analyst teams. Providing on-fly solutions to make database healthy and keep running during Change Control and Enterprise release. 24x7 production environment and provided on-call and day-to-day support. Environment: Oracle 10g/11g, Solaris 11, Red Hat Enterprise Linux 4, Cassandra 2.1, AWS EC2, Maria DB 10.1, DevCenter, Mongodb, Cqlsh, OpsCenter, Nodetool, ASM, HP-UX11, Shareplex, OEM (Grid Control), RMAN, RAC and Data Guard, Explain Plan, PL/SQL, Data Pump, Putty, SQLT, Sun Solaris 5.10. AT&T - Bedminister, NJ June 2005 to January 2011 AT&T Marketing Database (MDB) system, a massive database driven application used by sales force and others with feeds from and to many systems along with a web interface. Data analysis and business management based on

client specifications. Extensive use of SQL to write complex queries for reporting purpose. Extensively used PL/SQL to create Stored Procedures, Functions, Packages and Triggers. Involved in SQL*LOADER to load data from external files into tables in Oracle Database. Used IMP and EXP for loading and exporting tables. Performance tuning & SQL tuning using Explain Plan, TKPROF. Involved in Testing (Both Unit and Integration Level) and Implementation. Done the documentation and provided production support. Collecting user requirements and preparing the specifications for program units. Created simple, complex and materialized views for better performance and security. Created Test Data for testing procedures and functions. Debugging, editing and testing existing procedures and functions for enhanced functionality. Environment: Oracle 10g, TOAD, SQL Navigator 4.1.2, SQL, PL/SQL, Shell Scripts, Excel, Access, Cognos Sr. Oracle Developer Ernst & Young's LLP - Lyndhurst, NJ January 2004 to June 2005 Global Financial Information System Designed and created Batch Processing Job Scheduling. Created templates of Korn shell scripts for the various data feeds. Loaded data in staging, ODS, Data Marts of DW from PeopleSoft. Involved in 24/7 production support and implementing ICE Requests (Change Control). Created numerous Tables, Indexes, packages, procedures and functions for generic use. Optimizing SQL queries; used cost-based optimizer (CBO) hints to enhance Oracle performance. Worked on parallel execution of SQL statements and the direct-load Insert features. Used Dynamic SQL, as objects in data warehouse that are not pre-defined. Paralleling Inserts, Updates and Deletes by alter the session, enabling Parallel DML. Analyzed tables to estimate and compute statistics for the best execution plan. Created functional design and operational documentation. Frequent verification with Product Management and business customers of work in progress to ensure project delivered meets client's expectations. Utilized most of the Oracle 9i features, Materialize Views, Oracle table functions etc., Interacted with QA folks to help them write the unit test plan. Involved in Bench Marking Exercise in tuning the Batch Jobs. Environment: HP-UX B.11.11 U 9000/ 800, Windows 2000 Professional., Oracle9i Enterprise Edition Release 9.2.0.4.0, PL/ SQL Release 8.1.6, Korn Shell Scripting, SQL*Plus Release 8.1.6, SQL*Loader Release 8.1.6, TOAD 7.3, Vision Explorer 8.1, Informatica, Forms 6.0, Reports 6.0, Java Sr. Data Warehouse

Consultant Verizon Communications - Arlington, VA August 2001 to January 2004 Dimensional modeling to create the logical design using Erwin. Transferred the flat files into Oracle Database. Created various control files handling nulls and used direct path loading for performance. Created templates, korn shell scripts for the various data feeds. Created PL/SQL and SQL scripts for transformation of various data feeds. Used the best transformation techniques for key lookups using CTAS technique and etc. Used Dynamic SQL, as objects in data warehouse are not pre-defined (using latest feature of 8i "execute Immediate"). Paralleling Inserts, Updates and Deletes by alter the session, enabling Parallel DML. Aggregated Data Loads into Fact Tables from Fact Tables. Analyzed temporary tables to estimate and compute statistics for the best execution plan. Optimized SQL queries used cost-based optimizer (CBO) hints to enhance Oracle performance. Worked on parallel execution of SQL statements and the direct-load Insert features.

Created numerous packages, procedures and functions for generic use. Monitor the Hit Ratio of the database by using Dynamic Performance tables and usage of Tkprof, Utlbstat and Utlestat to identify and overcome performance Bottlenecks. Extensively worked on aggregations, largely used new family of analytic SQL functions such as rollup, cube etc. Involved as team player in Design of Data Warehouse/ Staging/ Mart, well versed with star schema, snowflake schema, type1, type2 and type 3 dimensions etc using ERwin. Involved in production support and implementing Change Control Requests (CCR's). Frequent verification with Product Management and business customers of work in progress to ensure project delivered met client's expectations. Created operational documentation. Environment: HP-UX B.11.11 U 9000/ 800, Windows NT version 4.0, Humming Bird V6.2, Erwin 3.5/ 4.0, Oracle8i Enterprise Edition 8.1.7, PL/ SQL Release 8.1.6, Korn Shell Scripting, SQL*Plus Release 8.1.6, SQL*Loader Release 8.1.6, TOAD 7.1.7, Informatica 5.1, Microstrategy 7. ETL Data Load/Senior Oracle Programmer AHoldUSA - Greenville, SC July 2000 to August 2001 Transferred flat files to TSM storage media. Extracted .dat and .mft flat files from TSM using ADSMPIPE. Created various control files handling nulls and used direct path loading for performance. Created and modified the templates, korn shell scripts for the various data feeds. Created PL/ SQL and SQL scripts for transformation of various data feeds. Used the best

transformation techniques for key lookups using CTAS technique and etc. Worked with star schema and well versed with 3NF and snowflake schema and optimized star queries by enabling the enabling the session with STAR_TRANSFORMATION_ENABLED to true. Created numerous temporary tables and indexes and populated parallel without logging highly optimizing storage parameters. Worked on aggregations, largely used new family of analytic SQL functions like rollup, cube etc. Used Dynamic SQL, as objects in data warehouse are not pre-defined (using latest feature of 8i "execute Immediate"). Paralleling Inserts, Updates and Deletes by alter the session, enabling Parallel DML. Aggregated Data Loads into Fact Tables from Fact Tables. Analyzed temporary tables to estimate and compute statistics for the best execution plan. Optimized SQL queries, used cost-based optimizer (CBO) hints to enhance oracle performance. Worked on parallel execution of SQL statements and the direct-load Insert features. Worked on gathering the statistics using DBMS_STATS package. Created numerous packages, procedures and functions for generic use. Environment: IBM/ AIX RISC System/ 6000 AIX 4.3, Windows NT version 4.0, Humming Bird V6.2, Oracle8i Enterprise Edition 8.1.6, PL/ SQL Release 8.1.6, SQL*Plus Release 8.1.6, SQL*Loader Release 8.1.6, SQL Navigator Version 3.2d11 (PVCS & MS VSS) Tuning Xpert, SQL Tools v1.4.1. Education Bachelor of Engineering in Technology and Science Warangal Engineering Kakatiya Institute of Technology 1995 Skills Jenkins, AWS, Linux, Devops Links <http://www.linkedin.com/in/srinivas-putta-57280415>

Name: Tina Martinez

Email: juliegraham@example.com

Phone: 884-679-0782