

Senior Consultant Senior Consultant Senior Consultant Austin, TX Work Experience Senior Consultant Avalon Consulting LLC 2014 to 2015 Staff Augmentation to Cloudera and Hortonworks Professional Services. * HDP security implementation for Telecom and Retail clients: Knox, Ranger, LDAP, and Kerberos integration. * CDH installation and upgrade for Manufacturing and Financial clients, including security reviews, Cloudera Navigator implementation, Kerberos and LDAP integration, HBase and Solr replication for disaster/recover - business continuity. * ElasticSearch/Kibana development of a services kiosk for a Government client. * Couchbase development in Node.JS, Javascript, Java 8 (Lambda) and Python. Contract Cloud Software Engineer Cisco Systems 2014 to 2014 Development of RESTful API middleware to interface Cisco Intelligent Automation for Cloud (IAC) client and backend Service Catalog/Process Orchestrator. * Integration of OpenStack, AWS, vCloudDirector, vSphere and Cisco UCS. * Development of a secure repository (keystore) for IAC client based on OpenStack's Barbican repository and Keystone authentication service. * Development of a generalized archiving/deletion/export facility for Cisco IAC 4.1. Contract Field Engineer EMC2 Corporation 2014 to 2014 Installation of Pivotal Hadoop distribution and Greenplum Database on EMC2 data appliance for two US clients. Client Services Engineer Zenoss, Inc 2013 to 2013 Developed Python extension modules (ZenPacks) to the Zenoss monitoring/administration system for Enterprise client-specific hardware and software components, including 3Par storage, CloudStack, Multiple Hadoop distributions, and Apache HBase. * Architecture Review of Open-software projects for next-generation Enterprise monitoring, including HBase, OpenTSDB, Memcache, and Spring Python Contract Senior Python Developer Hewlett-Packard 2012 to 2013 Sole responsibility for break/fix of RepDV (Reputation Digital Vaccine), a database of malicious IP addresses and DNS entries gathered worldwide, used to control HP TippingPoint Intrusion Protection System (IPS) and Security Monitoring System (SMS) appliances; re-establishment of a three-tier development/staging/production build environment; and re-architecture of RepDV to introduce greater modularity and isolation between data aggregation, database management, and package product production. * RepDV, which consists of Python and Twisted, maintains a 60GB database of exploit data that is graded, aged and completely turned over

every 24-hours through a batch update process run every 2 hours. Tipping Point IPS and SMS appliances worldwide are updated with deltas every 2 hours and full replacement packages every 12-24 hours, depending on model, as needed by threat severity. RepDV constitutes ~20% of HP's Enterprise Security Software revenue. * Evaluation of Hadoop for next-generation RepDV processing. Senior Architect A Cloud For Everyone 2011 to 2013 Designed and implemented a POC 10-node Apache Hadoop/CentOS cluster. * Designed a portfolio of Big Data/HPC clusters from 20 CPUs/112TB mini-rack (table top) to 144 CPU/0.5PB single rack costing \$50k to \$272k and marketed in the San Antonio, TX area. * Adding at-rest encryption to Apache Accumulo. Contract Senior VMware Virtualization Engineer US Army Medical Information Technology Center - Fort Sam Houston, TX 2011 to 2012 AD/Exchange Group, Systems Operations Branch, Sustainment Division, USAMITC * Tier 3 support of VMware Virtualization for Army Medical Command and Military Health System (MHS) users worldwide. * Upgrade of Datacenters from VMware Infrastructure 3 to vSphere 4/ESX to ESXi. * Virtualization of Microsoft Tier I applications (2008 SP 2/R2, AD, SharePoint, Exchange) including Physical to Virtual and direct Virtual. * Automation of VMware operations using PowerShell/PowerCLI on Windows and Python and Perl on Linux, including build out, configuration and management/reporting. * Maintenance of virtual networking at the L2 and L3 layer, VLANs and NIC bonding/teaming. * PKI-E System Administrator - responsible for the procurement and management of Army X.509 certificates for OSs, VMs and Application Servers for the AD/Exchange Group. * Build, Provision and Manage HP ProLiant Blade Servers, C7000 Enclosures and P4000 Storage systems. Contract System Administrator NewTek - San Antonio, TX 2011 to 2011 Designed and Implemented Zenoss and Nagios network operations/monitoring center. * SQL Server to Oracle migration of Enterprise CMS. * Security incident investigation. Owner / Senior Officer A Computer For Me - San Antonio, TX 1998 to 2010 A sole proprietorship software and hardware company dedicated to serving the requirements of the client. Specializing in secure Linux-based cluster computers for High-availability database, High-performance computational requirements, or Big Data/Batch processing * Built and operated a 6-node POC Hadoop cluster and developed MapReduce programs for a US Climate study that processed the entire NOAA

National Climatic Data Center temperature data set (~37GB) for 1901-2010 to determine decade-to-decade temperature change * Designed and operated Data Center for USAF clients on a 2-year contract, including Internet connection and local network, rack space, power, and cooling to support a surface chemistry study * Installed, stress-tested, and operated client specified software including operating systems (Linux), compilers (Portland Group HP FORTRAN and C/C++; Intel C++), databases (Oracle w/RAC, MySQL, Firebird embedded), web servers (Apache, IIF), parallel processing software (MPI-CH, PVM), parallel file systems (Luster; Red Hat GFS, monitoring software (Nagios); schedulers/queuing software (Maui scheduler; Moab; NQS) Data Layer Expert Trinity Millennium - San Antonio, TX 2007 to 2007 Implemented Python/C/MySQL replication for government mainframe re-engineering project. * Designed and implemented ETL using DTS and SSIS on SQL Server. Senior Research Analyst/Database Administrator Division 20 CNWRA Southwest Research Institute - San Antonio, TX 1988 to 1998 Responsible for the integrity and availability of data and systems for the Division, while providing all Information Systems Security functions from planning through auditing. * Implemented multiple mainframe Information Management Systems using AS, SQL/DS, REXX, RXSQL, and STAIRS and re-implemented mainframe systems and new client-server using GALAXY, Verity, Oracle, C++, PL/SQL and embedded SQL * Operated a Production/Test/Development environment by migrating changes as developed. * Tuned databases and optimized queries and indices * Architected of a high-level secure capacity network design, installed bastion host on DMZ with application gateway and implemented packet filter firewall on Cisco routers. Education Medicine University of Texas Medical School - San Antonio, TX 1980 to 1984 B.S. in Chemistry University of Notre Dame - Notre Dame, IN May 1980 Skills Information Security (10+ years), Hadoop Ecosystem (4 years), RDBMS (Oracle/MySQL/PostgreSQL/DB2) (10+ years), Linux (RHEL/CentOS/Ubuntu) / Unix / BDS / SunOS/Solaris (10+ years) Links <https://www.linkedin.com/in/robertmarshallcissp> <http://about.me/robert/marshall> <http://vizualize.me/robertlmarshall> <http://vizualize.me/robertlmarshall> <http://vizualize.me/robertlmarshall> Additional Information * Over [20] years of experience in full life-cycle design, development, and support of systems application architecture * Extensive

experience performing as lead database administrator on multiple mainframe and mini systems and multiple database implementations (relational, object, full text and NoSQL) * Extensive experience and training with security and information standards including HIPAA, FISMA, DISA STIGs and NIST * Over [10] years of hands-on experience designing, building, installing, administering and upgrading cluster systems (10-256 nodes) for high-performance, high-availability and "Big-Data" requirements * Designed and implemented Big Data infrastructure using multiple Hadoop distributions and subcomponents

Current Technical Skills

- * Virtualization: VMware, Citrix/Xen, KVM, VirtualBox, Docker
- * Operating Systems: Red Hat Enterprise Linux/CentOS; Debian/Ubuntu/Kubuntu; Windows 7, 8, 2008R2
- * Relational DBMS: SQL/DS; DB2; Oracle RAC; MySQL; PostgreSQL
- * Object DBMS: CACHE; SQLAlchemy
- * Hadoop Distributions: Stock Apache; Cloudera; Hortonworks; MapR M; Pivotal (Greenplum)
- * MapReduce/Hadoop Ecosystem: Hadoop, HDFS, HBase, Sqoop, Flume, Pig, Streams, Disco, Octopus, Cascading, Storm, Spark
- * Hadoop Security: Kerberos, LDAP, Knox, Ranger, Cloudera Navigator, HDFS Encryption Zones
- * NoSQL: HBase, Accumulo, Riak, Neo4J, Cassandra, MongoDB, Vertica, Couchbase
- * Languages: Python, Java, Javascript, Scala, Ruby, R, C/C++, bash
- * Monitoring: Zenoss; Nagios; Ganglia; Cacti; OpenTSDB
- * Messaging: Kafka, RabbitMQ
- * Version Control: Git, SVN, Mercurial
- * Provisioning: Vagrant, Puppet, Ansible, Ambari, Whirr, Mesos, Star Cluster
- * Development Frameworks: Griffon, Spring Data, Django, Flask
- * Cloud: vCloud Director, OpenStack, Cisco Intelligent Automation for Cloud (IAC)
- * AWS: StarCluster, Eucalyptus, Whirr
- * Text-Search: Lucene, Solr, Elasticsearch
- * Data Visualization: D3, IPython, NumPy, pandas, matplotlib, SciPy
- * Mobile: jQuery and jQuery mobile, Sencha Touch
- * Performance: Cython, LLVM
- * DevOps Tools: Jenkins, Confluence, JIRA, Github, Maven, Gradle

Name: Amber Diaz

Email: monicaharris@example.com

Phone: 269-586-2591